

MonsterKong - Agent - reinforcement learning

Alexandru Lucian - IIIA4

Onesim Robert - IIIA4

1. Etape:

- **Input: State-ul** - GameScreenRGB (getScreenRGB) **vector de pixeli**
- **Preprocesarea imaginilor**
- Calcularea functiei de scor (**reward-ul**) bazat pe state-ul jocului
- Introducerea imaginilor in **reteaua neuronală** (Q-learning algorithm - maximizarea reward-ul)
- **Output:** cea mai buna mutare (sus/jos/stanga/dreapta/saritura)

2. Input

Am folosit PLE: A Reinforcement Learning Env (<http://pygame-learning-environment.readthedocs.io/en/latest/>) pentru a simula jocul si a primi ca input state-ul acestuia.

3. Preprocesarea Imaginilor

- Transformam imaginile color primite in alb-negru
- Redimensionam imaginile in 80x80 pixeli
- Grupam cate patru imagini pentru ca modelul sa poata deduce viteza de miscare bilelor de foc si a monstruletului.

4. Reteaua Neuronala

Am folosit **Keras** cu **TensorFlow** backend

- 1 input layer - 4x80x80 image
- 3 convolution layers (hidden)
- 1 hidden fully connected layer
- 1 output fully connected layer- 5 noduri, cate unul pentru fiecare actiune

Functia de activare folosita ReLu ($f(x) = \max(0, x)$) - x input of a neuron

Distributie normala cu deviatia standard 0.01

Am folosit ADAM ([Adaptive Moment Estimation](#)) pentru optimizarea retelei

5. Flow-ul antrenamentului

- vom avea un numar de frame-uri (EXPLORE), perioada in care vom construi initial baza de antrenament (deoarece la inceput nu avem date)
- un numar de iteratii in care se agentul se va antrena (TRAIN)
- lista (replay memory) la care se adauga constant frame-uri (si in faza de explorare dar si in faza de training). Dimensiunea maxima este 20000 urmand ca pe parcurs (cand se umple) sa stergem cele mai vechi frame-uri. Din replay memory sunt alese batch-urile pentru training.
- epsilon (gradul de explorare) - parametru care va fi initial o valoare intre 0-1.
 - La fiecare pas al iteratiei se alege o valoare random intre 0-1:
 - daca valoarea e mai mica ca epsilon , se va alege o actiune random astfel oferindu-se sansa agentului de a explora state-uri noi.
 - daca valoarea este mare decat epsilon atunci modelul va prezice o actiune pe baza state-ului curent `model.predict(state)`;
- cu actiunea respectiva generam reward-ul si formam tuplu (s, a, r, s') care se adauga la replace memory
- in faza de training se alege un batch din replace memory de dimensiune 20
 - $Q(s,a) = r + \gamma \cdot \text{argmax}_{a'} Q(s', a')$ reward-ul imediat + o valoare maxima viitoare rezultata dintr-o stare viitoare
 - GAMMA = variabila de discount intre 0 si 1. Cu cat reward-ul e mai indepartat (mai viitor) cu atat mai putin il luam in considerare. Daca e 0, ne bazam pe reward-ul imediat.
 - `targets = model.predict(old_state)`
 - `newQval = model.predict(new_state)`
 - `target[moves] = reward + gamma * max(newQval)`.
 - `model.train_on_batch (inputs, targets)`

Referinte importante:

http://www.cdf.toronto.edu/~csc2542h/fall/material/csc2542f16_dqn.pdf

<http://files.davidqiu.com/research/nature14236.pdf>

<http://sebastianruder.com/optimizing-gradient-descent/>

<https://adeshpande3.github.io/adeshpande3.github.io/A-Beginner's-Guide-To-Understanding-Convolutional-Neural-Networks/>

| | | | | | | | |
|-------|-------|---------|---------------------|--------|------|--------|---------------------|
| FRAME | 60062 | EPSILON | 0.14313899999994315 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60063 | EPSILON | 0.14313799999994314 | ACTION | JUMP | REWARD | -6.215277777777778 |
| FRAME | 60064 | EPSILON | 0.14313699999994314 | ACTION | JUMP | REWARD | -1.2486111111111111 |
| FRAME | 60065 | EPSILON | 0.14313599999994314 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60066 | EPSILON | 0.14313499999994314 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60067 | EPSILON | 0.14313399999994314 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60068 | EPSILON | 0.14313299999994314 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60069 | EPSILON | 0.14313199999994314 | ACTION | LEFT | REWARD | 1.284722222222222 |
| FRAME | 60070 | EPSILON | 0.14313099999994314 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60071 | EPSILON | 0.14312999999994314 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60072 | EPSILON | 0.14312899999994314 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60073 | EPSILON | 0.14312799999994313 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60074 | EPSILON | 0.14312699999994313 | ACTION | LEFT | REWARD | 1.284722222222222 |
| FRAME | 60075 | EPSILON | 0.14312599999994313 | ACTION | DOWN | REWARD | 1.3541666666666665 |
| FRAME | 60076 | EPSILON | 0.14312499999994313 | ACTION | LEFT | REWARD | 1.284722222222222 |
| FRAME | 60077 | EPSILON | 0.14312399999994313 | ACTION | DOWN | REWARD | 1.3541666666666665 |
| FRAME | 60078 | EPSILON | 0.14312299999994313 | ACTION | JUMP | REWARD | 1.3541666666666665 |
| FRAME | 60079 | EPSILON | 0.14312199999994313 | ACTION | JUMP | REWARD | 1.3541666666666665 |
| FRAME | 60080 | EPSILON | 0.14312099999994313 | ACTION | JUMP | REWARD | 1.3541666666666665 |
| FRAME | 60081 | EPSILON | 0.14311999999994313 | ACTION | JUMP | REWARD | 1.3541666666666665 |
| FRAME | 60082 | EPSILON | 0.14311899999994313 | ACTION | JUMP | REWARD | -6.145833333333334 |
| FRAME | 60083 | EPSILON | 0.14311799999994312 | ACTION | JUMP | REWARD | -1.1791666666666667 |
| FRAME | 60084 | EPSILON | 0.14311699999994312 | ACTION | JUMP | REWARD | 1.3541666666666665 |
| FRAME | 60085 | EPSILON | 0.14311599999994312 | ACTION | JUMP | REWARD | 1.3541666666666665 |
| FRAME | 60086 | EPSILON | 0.14311499999994312 | ACTION | JUMP | REWARD | 1.3541666666666665 |
| FRAME | 60087 | EPSILON | 0.14311399999994312 | ACTION | JUMP | REWARD | 1.3541666666666665 |
| FRAME | 60088 | EPSILON | 0.14311299999994312 | ACTION | NOOP | REWARD | 1.3541666666666665 |
| FRAME | 60089 | EPSILON | 0.14311199999994312 | ACTION | JUMP | REWARD | 1.3541666666666665 |
| FRAME | 60090 | EPSILON | 0.14311099999994312 | ACTION | LEFT | REWARD | 1.284722222222222 |
| FRAME | 60091 | EPSILON | 0.14310999999994312 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60092 | EPSILON | 0.14310899999994312 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60093 | EPSILON | 0.14310799999994311 | ACTION | LEFT | REWARD | 1.284722222222222 |
| FRAME | 60094 | EPSILON | 0.1431069999999431 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60095 | EPSILON | 0.1431059999999431 | ACTION | DOWN | REWARD | 1.3541666666666665 |
| FRAME | 60096 | EPSILON | 0.1431049999999431 | ACTION | LEFT | REWARD | 1.284722222222222 |
| FRAME | 60097 | EPSILON | 0.1431039999999431 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60098 | EPSILON | 0.1431029999999431 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60099 | EPSILON | 0.1431019999999431 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60100 | EPSILON | 0.1431009999999431 | ACTION | JUMP | REWARD | 1.284722222222222 |
| FRAME | 60101 | EPSILON | 0.1430999999999431 | ACTION | JUMP | REWARD | -6.215277777777778 |
| FRAME | 60102 | EPSILON | 0.1430989999999431 | ACTION | JUMP | REWARD | -1.2486111111111111 |
| FRAME | 60103 | EPSILON | 0.1430979999999431 | ACTION | UP | REWARD | -1.2486111111111111 |
| FRAME | 60104 | EPSILON | 0.1430969999999431 | ACTION | DOWN | REWARD | -1.1791666666666667 |

```

231         player.startNetwork(99999999, player.EPSILON_FINAL)
232
233     elif sys.t
234         player
235
236     else:
237         print
238
239     if name ==
240         main()

```

pygame window

```

named
FRAME 2377 EPSILON 0.0001 ACTION JUMP REWARD 3.4333333333333336
FRAME 2378 EPSILON 0.0001 ACTION UP REWARD 3.6
FRAME 2379 EPSILON 0.0001 ACTION UP REWARD 3.7666666666666666
FRAME 2380 EPSILON 0.0001 ACTION RIGHT REWARD 3.6
FRAME 2381 EPSILON 0.0001 ACTION NOOP REWARD 3.6
FRAME 2382 EPSILON 0.0001 ACTION NOOP REWARD 3.6
FRAME 2383 EPSILON 0.0001 ACTION RIGHT REWARD 3.4333333333333336
FRAME 2384 EPSILON 0.0001 ACTION RIGHT REWARD 3.2666666666666666
FRAME 2385 EPSILON 0.0001 ACTION JUMP REWARD 3.2666666666666666
FRAME 2386 EPSILON 0.0001 ACTION RIGHT REWARD 3.1
FRAME 2387 EPSILON 0.0001 ACTION RIGHT REWARD 2.9333333333333336
FRAME 2388 EPSILON 0.0001 ACTION JUMP REWARD 2.9333333333333336
FRAME 2389 EPSILON 0.0001 ACTION DOWN REWARD 2.9333333333333336
FRAME 2390 EPSILON 0.0001 ACTION LEFT REWARD 2.9333333333333336
FRAME 2391 EPSILON 0.0001 ACTION UP REWARD 3.1
FRAME 2392 EPSILON 0.0001 ACTION DOWN REWARD 3.1
FRAME 2393 EPSILON 0.0001 ACTION LEFT REWARD 3.1
FRAME 2394 EPSILON 0.0001 ACTION JUMP REWARD 3.1

```

```

COIN_WEIGHT_Y = 5.
COIN_VALUE = 3.0 # the initial weight of coin proximity

ACTIONS = 5 # valid moves
EPSILON_INITIAL = 0.2 # starting value of epsilon
EPSILON_FINAL = 0.0001 # final value of epsilon

def __init__(self):
    self.buildModel()
    self.game = MonsterKong()
    self.p = PLE(self.game, fps=30, display_screen=True)

def buildModel(self):
    # build model using Keras
    self.model = Sequential()
    # input layer 4x80x80 image and 3 convolution hidden layers with activation function "ReLu" f(x
    self.model.add(Convolution2D(32, 8, 8, subsample=(4, 4), init=lambda shape, name: normal(shape,
    self.model.add(Activation('relu'))
    self.model.add(Convolution2D(64, 4, 4, subsample=(2, 2), init=lambda shape, name: normal(shape,
    self.model.add(Activation('relu'))
    self.model.add(Convolution2D(64, 3, 3, subsample=(1, 1), init=lambda shape, name: normal(shape,
    self.model.add(Activation('relu'))
    self.model.add(Flatten())
    # last hidden fully connected layer
    self.model.add(Dense(512, init=lambda shape, name: normal(shape, scale=0.01, name=name)))
    self.model.add(Activation('relu'))
    # output layer - 1 neuron for each valid move (5)
    self.model.add(Dense(5, init=lambda shape, name: normal(shape, scale=0.01, name=name)))

    # using ADAM (Adaptive Moment Estimation)
    self.model.compile(loss='mse', optimizer=Adam(lr=1e-6))

def saveModel(self):

```